Autonomous tool creation, execution and reuse

Search and test evidence shown separately · benchmark / built-in tools excluded

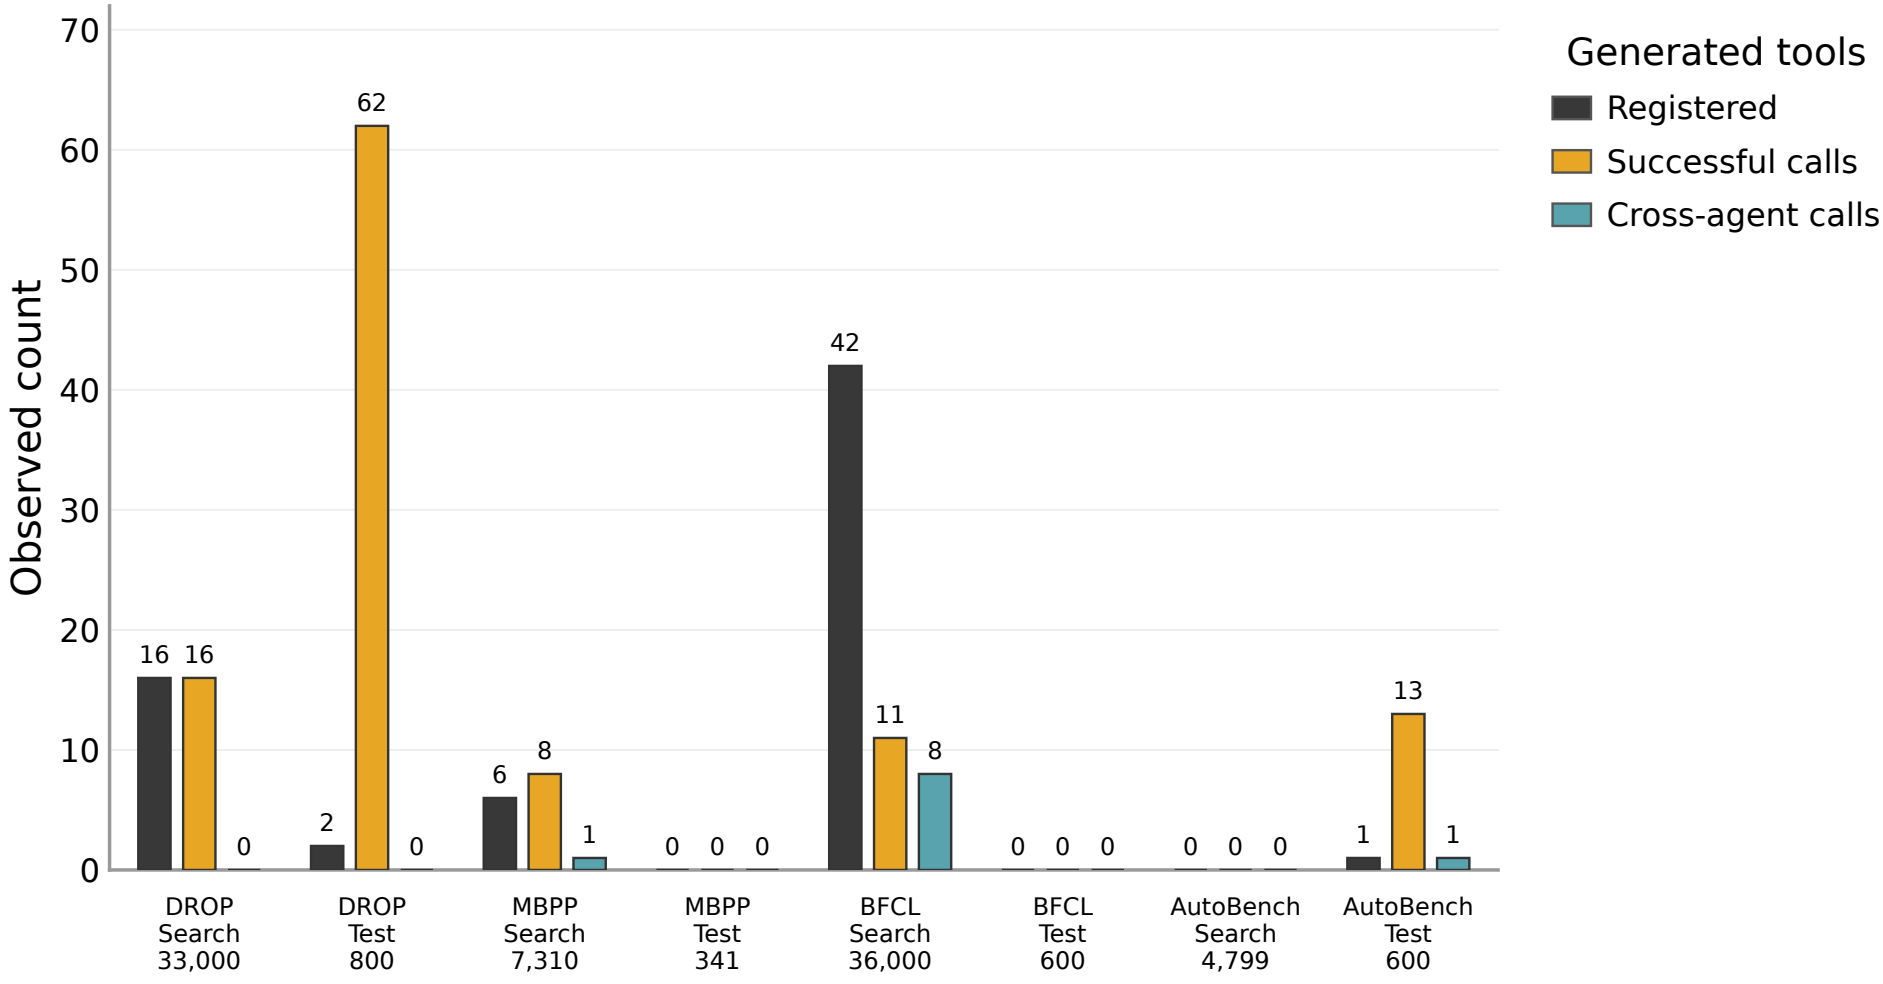

Numbers below stage labels are observed episodes. Search = all retained task × pass observations; test = unique tasks.

Exposure differs across stages; raw counts describe observed use, not success rates. Cross-agent calls are a subset of successful calls.

Available tool traces only; coverage is in tools.csv. MATH is unknown; HLE has no full test. AutoBench = AutomationBench v11.